

1 Overview

In this project, you will implement a multiprocessor operating system simulator using a popular userspace threading library for linux called pthreads. The framework for the multithreaded OS simulator is nearly complete, but missing one critical component: the CPU scheduler! Your task is to implement the CPU scheduler, using three different scheduling algorithms.

Note: Make sure that multiple CPU cores are enabled in your virtual machine, otherwise you will receive incorrect results. See the TAs if you need help.

This can be done in VirtualBox by going to: Settings > System > Processor > CPUs

If you are using the Vagrant setup, the number of cores should default to 2. You can run `nproc --all` to see how many cores are available to your VM.

We have provided you with source files that constitute the framework for your simulator. You will only need to modify `answers.txt` and `student.c`. However, just because you are only modifying two files doesn't mean that you should ignore the other ones - there is helpful information in the other files. Information about using the pthreads library is given in Problem 0. We have provided you these files:

1. `Makefile` - Working one provided for you; do not modify.
2. `os-sim.c` - Code for the operating system simulator which calls your CPU scheduler.
3. `os-sim.h` - Header file for the simulator.
4. `process.c` - Descriptions of the simulated processes.
5. `process.h` - Header file for the process data.
6. `student.c` - This file contains stub functions for your CPU scheduler.
7. `student.h` - Header file for your code to interface with the OS simulator

1.1 Scheduling Algorithms

For your simulator, you will implement the following three CPU scheduling algorithms:

1. First In, First Out (FIFO) - Runnable processes are kept in a ready queue. FIFO is non-preemptive; once a process begins running on a CPU, it will continue running until it either completes or blocks for I/O.
2. Round-Robin - Similar to FIFO, except preemptive. Each process is assigned a timeslice when it is scheduled. At the end of the timeslice, if the process is still running, the process is preempted, and moved to the tail of the ready queue.
3. Shortest Remaining Time First (SRTF) - The process with the shortest remaining time in its burst always gets the CPU. Longer processes must be pre-empted if a process that has a shorter burst becomes runnable.

1.2 Process States

In our OS simulation, there are five possible states for a process, which are listed in the process state `enum` in `os-sim.h`:

1. NEW - The process is being created, and has not yet begun executing.
2. READY - The process is ready to execute, and is waiting to be scheduled on a CPU.
3. RUNNING - The process is currently executing on a CPU.

4. WAITING - The process has temporarily stopped executing, and is waiting on an I/O request to complete.

5. TERMINATED - The process has completed.

There is a field named state in the PCB, which must be updated with the current state of the process. The simulator will use this field to collect statistics.

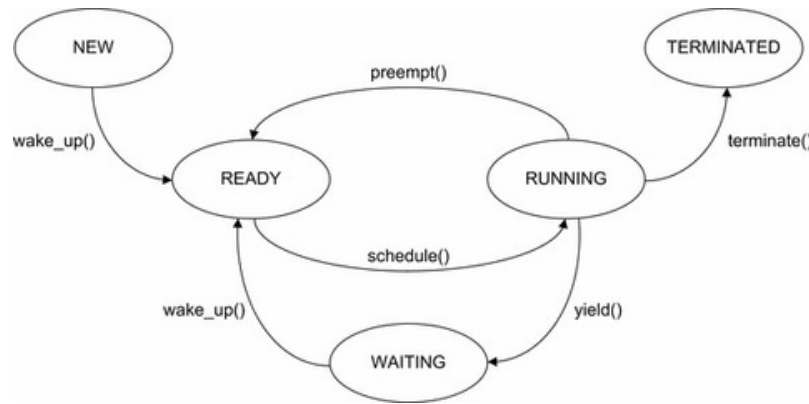


Figure 1: Process States

1.3 The Ready Queue

On most systems, there are a large number of processes, but only one or two CPUs on which to execute them. When there are more processes ready to execute than CPUs, processes must wait in the READY state until a CPU becomes available. To keep track of the processes waiting to execute, we keep a ready queue of the processes in the READY state

Since the ready queue is accessed by multiple processors, which may add and remove processes from the ready queue, the ready queue must be protected by some form of synchronization—for this project, you will use a mutex lock. The ready queue SHOULD use a different mutex than the current mutex.

1.4 Scheduling Processes

schedule() is the core function of the CPU scheduler. It is invoked whenever a CPU becomes available for running a process. schedule() must search the ready queue, select a runnable process, and call the context switch() function to switch the process onto the CPU.

There is a special process, the idle process, which is scheduled whenever there are no processes in the READY state.

1.5 CPU Scheduler Invocation

There are four events which will cause the simulator to invoke schedule():

1. yield() - A process completes its CPU operations and yields the processor to perform an I/O request.
2. wake up() - A process that previously yielded completes its I/O request, and is ready to perform CPU operations. wake up() is also called when a process in the NEW state becomes runnable.
3. preempt() - When using a Round-Robin or SRTF scheduling algorithm, a CPU-bound process may be preempted before it completes its CPU operations.

Time	Ru	Re	Wa	CPU 0	CPU 1	< I/O Queue >
=====	==	==	==	=====	=====	=====
0.0	0	0	0	(IDLE)	(IDLE)	< <
0.1	0	0	0	(IDLE)	(IDLE)	< <
0.2	0	0	0	(IDLE)	(IDLE)	< <
0.3	0	0	0	(IDLE)	(IDLE)	< <
0.4	0	0	0	(IDLE)	(IDLE)	< <
0.5	0	0	0	(IDLE)	(IDLE)	< <
0.6	0	0	0	(IDLE)	(IDLE)	< <
0.7	0	0	0	(IDLE)	(IDLE)	< <
0.8	0	0	0	(IDLE)	(IDLE)	< <
0.9	0	0	0	(IDLE)	(IDLE)	< <
1.0	0	0	0	(IDLE)	(IDLE)	< <
.....						

Figure 3: Sample Output

The simulator generates a Gantt Chart, showing the current state of the OS at every 100ms interval. The leftmost column shows the current time, in seconds. The next three columns show the number of Running, Ready, and Waiting processes, respectively. The next two columns show the process currently running on each CPU. The rightmost column shows the processes which are currently in the I/O queue, with the head of the queue on the left and the tail of the queue on the right.

As you can see, nothing is executing. This is because we have no CPU scheduler to select processes to execute! Once you complete Problem 1 and implement a basic FIFO scheduler, you will see the processes executing on the CPUs.

2 Problem 1: FIFO Scheduler

NOTE: Part B of each part requires you to put your answer down in answers.txt

Part A. Implement the CPU scheduler using the FIFO scheduling algorithm. You may do this however you like, however, we suggest the following:

- Implement a thread-safe ready queue using a linked list. A linked list will allow you to reuse this ready queue for the Round-Robin and SRTF scheduling algorithms.
- Implement the `yield()`, `wake up()`, and `terminate()` handlers. `preempt()` is not necessary for this stage of the project. See the overview and the comments in the code for the proper behavior of these events.
- Implement `idle()`. `idle()` must wait on a condition variable that is signalled whenever a process is added to the ready queue.
- Implement `schedule()`. `schedule()` should extract the first process in the ready queue, then call `context switch()` to select the process to execute. If there are no runnable processes, `schedule()` should call `context switch()` with a NULL pointer as the PCB to execute the idle process.

2.1 Hints

- Be sure to update the state field of the PCB. The library will read this field to generate the Running, Ready, and Waiting columns, and to generate the statistics at the end of the simulation.
 - There is a field in the PCB, `next`, which you may use to build linked lists of PCBs.
 - Four of the five entry points into the scheduler (`idle()`, `yield()`, `terminate()`, and `preempt()`) should cause a new process to be scheduled on the CPU. In your handlers, be sure to call `schedule()`, which will select a runnable process, and then call `context switch()`. When these four functions return, the library will simulate the process selected by `context switch()`.
 - `context switch()` takes a timeslice parameter, which is used for preemptive scheduling algorithms. Since FIFO is non-preemptive, use -1 for this parameter to give the process an infinite timeslice.
- Part B. Run your OS simulation with 1, 2, and 4 CPUs. Compare the total execution time of each. Is there a linear relationship between the number of CPUs and total execution time? Why or why not? Keep in mind that the execution time refers to the simulated execution time.

3 Problem 2: Round-Robin Scheduler

Part A. Add Round-Robin scheduling functionality to your code. You should modify `main()` to add a command line option, `-r`, which selects the Round-Robin scheduling algorithm, and accepts a parameter, the length of the timeslice. For this project, timeslices are measured in tenths of seconds. E.g.:

```
./os-sim <# CPUs> -r 5
```

should run a Round-Robin scheduler with timeslices of 500 ms. While:

```
./os-sim <# of CPUs>
```

should continue to run a FIFO scheduler. You should also make sure `preempt` is implemented in this section of the project.

To specify a timeslice when scheduling a process, use the `timeslice` parameter of `context switch()`. The simulator will automatically preempt the process and call your `preempt()` handler if the process executes on the CPU for the length of the timeslice without terminating or yielding for I/O.

Part B. Run your Round-Robin scheduler with timeslices of 800ms, 600ms, 400ms, and 200ms. Use only one CPU for your tests. Compare the statistics at the end of the simulation. Show that the total waiting time decreases with shorter timeslices. However, in a real OS, the shortest timeslice possible is usually not the best choice. Why not?

4 Problem 3: Shortest Remaining Time First Scheduler

Part A. Add SRTF scheduling to your code. Modify `main()` to accept the `-s` parameter to select the SRTF algorithm. The `-r` and default FIFO scheduler should continue to work.

The scheduler should use the `time remaining` field of the PCB to prioritize processes that have a shorter remaining time in their CPU burst.

For SRTF scheduling, you will need to make use of the `current[]` array and `force preempt()` function. The `current[]` array should be used to keep track of the process currently executing on each CPU. Since this array is accessed by multiple CPU threads, it must be protected by a mutex. `current mutex` has been provided for you.

The `force_preempt()` function preempts a running process before its timeslice expires. Your `wake_up()` handler should make use of this function to preempt a process when a process with lower time remaining needs a CPU.

Part B. While it is easy to simulate an SRTF algorithm in the simulator, it is essentially impossible to implement precisely in real life and is thus usually approximated. Why is this the case?

Run each of the scheduling algorithms using one CPU and compare the total waiting times. Which one had the lowest? Why?

5 Deliverables

NOTE: Each Problem has two parts (labeled A and B). The first is the actual implementation, and the second is a question linked to the scheduling algorithm you are implementing. Make sure you complete both.

You can run `make submit` to automatically package your project for submission. Submit the resulting `tar.gz` compressed folder on T-square.

The archive should contain the following files:

- `answers.txt` - Short answers to questions from above.
- `Makefile` - Working one provided for you; don't break it.
- `os-sim.c` - Code for the operating system simulator.
- `os-sim.h` - Header file for the simulator.
- `process.c` - Descriptions of the simulated processes.
- `process.h` - Header file for the process data.
- `student.c` - Your code for the scheduler.
- `student.h` - Header file for your scheduler code.